

Hadoop Average temperature

1BM22CS066

```
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
localhost: namenode is running as process 5290. Stop it first and ensure /tmp/hadoop-hadoop-namenode.pid file is empty before retry.
Starting datanodes
localhost: datanode is running as process 5435. Stop it first and ensure /tmp/hadoop-hadoop-datanode.pid file is empty before retry.
Starting secondary namenodes [bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC]
bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC: secondarynamenode is running as process 5699. Stop it first and ensure /tmp/hadoop-hadoop-secondarynamenode.pid file is empty before retry.
Starting resourcemanager
resourcemanager is running as process 6069. Stop it first and ensure /tmp/hadoop-hadoop-resourcemanager.pid file is empty before retry.
Starting nodemanagers
localhost: nodemanager is running as process 6237. Stop it first and ensure /tmp/hadoop-hadoop-nodemanager.pid file is empty before retry.
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$
```

```
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
localhost: namenode is running as process 5290. Stop it first and ensure /tmp/hadoop-hadoop-namenode.pid file is empty before retry.
Starting datanodes
localhost: datanode is running as process 5435. Stop it first and ensure /tmp/hadoop-hadoop-datanode.pid file is empty before retry.
Starting secondary namenodes [bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC]
bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC: secondarynamenode is running as process 5699. Stop it first and ensure /tmp/hadoop-hadoop-secondarynamenode.pid file is empty before retry.
Starting resourcemanager
resourcemanager is running as process 6069. Stop it first and ensure /tmp/hadoop-hadoop-resourcemanager.pid file is empty before retry.
Starting nodemanagers
localhost: nodemanager is running as process 6237. Stop it first and ensure /tmp/hadoop-hadoop-nodemanager.pid file is empty before retry.
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -copyFromLocal /home/hadoop/Desktop/1901.txt /shuraim/temperature.txt
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -rm /shuraim/temperature.txt
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$
```

```

hadoop@bmsccese-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop jar /home/hadoop/Desktop/TemperatureAvg.jar AverageDriver /shurain/temperature.txt /shurain/tempOut
2025-05-06 15:22:59,024 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-05-06 15:22:59,064 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-05-06 15:22:59,065 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2025-05-06 15:22:59,127 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-05-06 15:22:59,179 INFO input.FileInputFormat: Total input files to process : 1
2025-05-06 15:22:59,207 INFO mapreduce.JobSubmitter: number of splits:1
2025-05-06 15:22:59,272 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1881755701_0001
2025-05-06 15:22:59,272 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-05-06 15:22:59,335 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-05-06 15:22:59,335 INFO mapreduce.Job: Running job: job_local1881755701_0001
2025-05-06 15:22:59,336 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-05-06 15:22:59,339 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 15:22:59,340 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:22:59,340 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:22:59,340 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2025-05-06 15:22:59,430 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-05-06 15:22:59,431 INFO mapred.LocalJobRunner: Starting task: attempt_local1881755701_0001_m_000000_0
2025-05-06 15:22:59,442 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 15:22:59,442 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:22:59,442 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:22:59,451 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2025-05-06 15:22:59,453 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/shurain/temperature.txt:0+888190
2025-05-06 15:22:59,488 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
2025-05-06 15:22:59,488 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2025-05-06 15:22:59,488 INFO mapred.MapTask: soft limit at 83886080
2025-05-06 15:22:59,488 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2025-05-06 15:22:59,488 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2025-05-06 15:22:59,490 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2025-05-06 15:22:59,575 INFO mapred.LocalJobRunner:
2025-05-06 15:22:59,576 INFO mapred.MapTask: Starting flush of map output
2025-05-06 15:22:59,576 INFO mapred.MapTask: Spilling map output
2025-05-06 15:22:59,576 INFO mapred.MapTask: bufstart = 0; bufend = 59076; bufvoid = 104857600
2025-05-06 15:22:59,576 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26188144(104752576); length = 26253/6553600
2025-05-06 15:22:59,585 INFO mapred.MapTask: Finished spill 0
2025-05-06 15:22:59,591 INFO mapred.Task: Task:attempt_local1881755701_0001_m_000000_0 is done. And is in the process of committing
2025-05-06 15:22:59,593 INFO mapred.LocalJobRunner: map
2025-05-06 15:22:59,593 INFO mapred.Task: Task 'attempt_local1881755701_0001_m_000000_0' done.
2025-05-06 15:22:59,596 INFO mapred.Task: Final Counters for attempt_local1881755701_0001_m_000000_0: Counters: 23
File System Counters
FILE: Number of bytes read=4258
FILE: Number of bytes written=716898
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=888190
HDFS: Number of bytes written=0
HDFS: Number of read operations=5
HDFS: Number of large read operations=0
HDFS: Number of write operations=1
HDFS: Number of bytes read erasure-coded=0

```

Hadoop Mean Max temperature

1BM22CS066


```

hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /shurain/tempOut
Found 2 items
-rw-r--r-- 1 hadoop supergroup          0 2025-05-06 15:22 /shurain/tempOut/_SUCCESS
-rw-r--r-- 1 hadoop supergroup          8 2025-05-06 15:22 /shurain/tempOut/part-r-00000
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cat /shurain/tempOut/part-r-00000
1901 46
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop jar /home/hadoop/Desktop/MeanMaxTemp.jar MeanMaxDriver /shurain/temperature.txt /shurain/tempOut1
Exception in thread "main" java.lang.ClassNotFoundException: MeanMaxDriver
    at java.base/java.net.URLClassLoader.findClass(URLClassLoader.java:476)
    at java.base/java.lang.ClassLoader.loadClass(ClassLoader.java:594)
    at java.base/java.lang.ClassLoader.loadClass(ClassLoader.java:527)
    at java.base/java.lang.Class.forName0(Native Method)
    at java.base/java.lang.Class.forName(Class.java:398)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:321)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:241)
hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop jar /home/hadoop/Desktop/MeanMaxTemp.jar MeanMaxDriver /shurain/temperature.txt /shurain/tempOut1
2025-05-06 15:24:42,211 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-05-06 15:24:42,254 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-05-06 15:24:42,254 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2025-05-06 15:24:42,315 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-05-06 15:24:42,368 INFO input.FileInputFormat: Total input files to process : 1
2025-05-06 15:24:42,396 INFO mapreduce.JobSubmitter: number of splits:1
2025-05-06 15:24:42,461 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local78876621_0001
2025-05-06 15:24:42,461 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-05-06 15:24:42,523 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-05-06 15:24:42,523 INFO mapreduce.Job: Running job: job_local78876621_0001
2025-05-06 15:24:42,524 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-05-06 15:24:42,527 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 15:24:42,528 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:24:42,528 INFO output.FileOutputCommitter: FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:24:42,528 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2025-05-06 15:24:42,578 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-05-06 15:24:42,578 INFO mapred.LocalJobRunner: Starting task: attempt_local78876621_0001_m_000000_0
2025-05-06 15:24:42,591 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 15:24:42,591 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:24:42,591 INFO output.FileOutputCommitter: FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:24:42,600 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2025-05-06 15:24:42,602 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/shurain/temperature.txt:0+888190
2025-05-06 15:24:42,638 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
2025-05-06 15:24:42,638 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2025-05-06 15:24:42,638 INFO mapred.MapTask: soft limit at 83886080
2025-05-06 15:24:42,638 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2025-05-06 15:24:42,638 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2025-05-06 15:24:42,640 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2025-05-06 15:24:42,717 INFO mapred.LocalJobRunner:

```

```
Map-Reduce Framework
  Bytes Read=888190
  Bytes Written=74
```

```
Map input records=6565
Map output records=6564
Map output bytes=45948
Map output materialized bytes=59082
Input split bytes=110
Combine input records=0
Combine output records=0
Reduce input groups=12
Reduce shuffle bytes=59082
Reduce input records=6564
Reduce output records=12
Spilled Records=13128
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=0
Total committed heap usage (bytes)=1052770304
```

```
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
```

```
File Input Format Counters
  Bytes Read=888190
File Output Format Counters
  Bytes Written=74
```

```
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /shuraim/tempOut1
Found 2 items
-rw-r--r-- 1 hadoop supergroup          0 2025-05-06 15:24 /shuraim/tempOut1/_SUCCESS
-rw-r--r-- 1 hadoop supergroup        74 2025-05-06 15:24 /shuraim/tempOut1/part-r-00000
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cat /shuraim/tempOut1/part-r-00000
```

```
01      4
02      0
03      7
04     44
05    100
06    168
07    219
08    198
09    141
10    100
11     19
12      3
```

```
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~$
```